



Fwd: Round 2 is open: Assembly. Submissions close Friday 25 September, 11:59 PM IST

1 message

Muskan Sharma <muskan.sharma2027@mastersunion.org>
To: Om Umrana <om.umrana2028@mastersunion.org>

30 September 2026 at 18:28

----- Forwarded message -----

From: Adhavan RK <adhavan@the-ken.com>

Date: Fri, 18 Sept 2026 at 21:21

Subject: Round 2 is open: Assembly. Submissions close Friday 25 September, 11:59 PM IST

To: <muskan.sharma2027@mastersunion.org>



Hi Muskan ,

Congratulations! You're through to Round 2, and one step closer to the finale, where things narrow and the prizes get real: a share of the ₹20 lakh prize pool, a ₹10 lakh grand prize at the top of it, plus a credit in the agentic map. Get through this round and you reach the Build round, where the field thins out and the end comes into view.

Round 1 asked you for an idea sketch. Round 2 turns that sketch into a design: an agent specified well enough that someone else could build it, and a clear statement of what the rails cannot do for you yet. The strongest designs go to the next round, where the rails open up and you find out whether your design survives contact with them.

This round, every team answers the same eight questions, whichever track you picked. The tracks split in the next round: Build teams go on to build their agent, and Strategy teams get a different problem statement.

What you submit

Your answers to the eight questions below, typed into the form. Plus, optionally, a link to your conversation with whatever AI tool you used.

Submission closes Friday, 25th September, 11:59 PM IST. The submission portal isn't live yet, we'll email you once it opens, so keep an eye on your inbox and the site.

The eight questions

- What is the outcome your agent is accountable for? One sentence.
- What is the level of autonomy your agent has? Score it L1 to L4 using this table. Your agent's level is set by the most consequential thing it does without asking anybody.

Level	What the agent does
L1	Prepares things for the person. They still do the work.
L2	Proposes a specific action every time. The person says yes or no to each one.
L3	Acts on its own inside limits somebody set. The person deals with whatever falls outside those limits.
L4	Plans several steps, does them, checks its own work, comes back when stuck, and is judged on whether the thing got done.

- What are the states your agent goes through? Share the happy flow and the unhappy flow. Be as specific as possible, and illustrate with a diagram if you prefer.
- On each of the three rails, what does your agent leverage that already exists, and what capabilities does it need that don't exist and must be built? Refer to the documentation for each rail:

- Gnani, voice: [documentation](#). Sign up on the [platform](#) as well, since the documentation alone will not show you everything it does. You'll get free credits courtesy Gnani to play with the platform.
- Pine Labs, payments and authorisation: pinelabs.com/docs
- Delhivery, logistics: [Developer portal](#) and [Maps reference](#)

For example, a capability table might look like:

Name	Used at	Rail	What you send it	What comes back	When it fails	What it must never do	What exists today	Build status
confirm_permission_status	State 3	Payments	customer reference, service identifier	whether it is active, the amount, how often, when it last ran	says so and stops, never assumes there is none	touch a permission the person has not named	[named as its documentation names it, and where you found it]	EXISTS

- If you had a choice, does your agent need a fourth rail to become much more effective? What would that rail do? Which Indian company would you want to build it?
- How will a human interact with your agent? What's their interface? Be as specific as possible.
- What is the name of your agent?
- Which Indian company has the best chance of creating an agent just like yours?

How this gets judged

As with every submission so far, judging for this round follows the same criteria as earlier rounds: evidence, creativity, clarity, feasibility and thoroughness.

To identify the strongest submissions with the most original work, we'll supplement human evaluation with the following:

- All submissions will be graded relative to other solutions within the same opening you've picked. So if you chose, say, *Madam OTP*, your solution will be compared for originality and insight against all other teams who picked that same space.
- All submissions will also be graded against our internal AI solution as a baseline. Every opening has been run through frontier AI tools, and your solution will be compared against those results.
- You can use any AI you like, but we ask teams to submit their logs and chats, which we'll use to rank and grade human versus machine contributions. We encourage the use of AI, but show your work.

FAQs

We cannot code. Are we at a disadvantage?

No. Nothing asked needs you to know how to code. All it requires is clear thinking about user behaviour and what a capability does. If the documentation of the rail partners feels intimidating, we encourage you to run it through an AI to understand it better.

We are engineers. Can we just build it?

Not yet. What gets scored is the design and what you found in the documentation.

Our agent needs something that does not exist. Do we lose marks?

No. In fact, the opposite — finding it, specifying it, and saying what your agent does in the meantime is one of the highest-scoring things you can do.

Do we have to build or test anything this round?

No. This round happens on paper. Next round is where submissions get separated between the Build track and the Strategy track, and teams on the Build track will be expected to build.

Where do we ask questions?

Write to adhavan@the-ken.com.

All the best,

Adhavan RK

This email was sent to muskan.sharma2027@mastersunion.org

Something wrong? Tell us at support@the-ken.com

Don't want to hear from us again? [Unsubscribe from our emails](#)

© 2026 The Ken